




# Randomized algorithms 1

## Intro, hashing

CS240

Spring 2022

*Rui Fan*



# Outline

- Introduction
- Probability review
- Max-cut and randomized quicksort
- Hashing
  - Closed addressing
  - Universal hashing
  - Perfect hashing


# Randomized algorithms

- Till now, all of our algorithms have been deterministic.
  - Given an input, the algorithm always does the same thing.
- It turns out it's very useful to allow algorithms to be nondeterministic.
  - As the algorithm operates, it's allowed to make some random choices.
  - Running the algorithm multiple times on same input can produce different behaviors.



# Why randomized algorithms?

- For many problems, randomized algorithms work better than deterministic ones.
  - Faster / uses less memory
  - Simpler, easier to understand.
  - Some problems that provably can't be solved (or solved efficiently) by deterministic algorithms can be solved by randomized ones.
  - According to quantum mechanics, the world is inherently probabilistic, so nature is randomized!






# How can randomness help?


- Say you have a string of length  $n$  that's half A's and half B's.
- We want to find a location in the string with an A.
- Any deterministic algorithm takes  $n/2+1$  steps in the worst case.
- But by checking random locations, a randomized algorithm finds an A in 2 steps in expectation.

# How can randomness help?

- Measure the area of this



- Method 1



- Method 2

- ❖ Print the shape out on a piece of paper.
- ❖ Throw 100 darts at it.
- ❖ See what percent land in the shape.
- ❖ Multiply by area of your paper.



# Las Vegas vs Monte Carlo

- A Las Vegas randomized algorithm always **produces the right answer**. But its running time can vary depending on its random choices.
  - We want to minimize the expected running time of a Las Vegas algorithm.
- A Monte Carlo algorithm always has the **same running time**. But it sometimes produces the wrong answer, depending on its random choices.
  - We want to minimize the error probability of a Monte Carlo algorithm.





# Probability review

- Discrete probability theory is based on events and their probabilities.
  - Events can be composed of more basic events.
  - Ex Event of rolling a 2 on a fair dice, with probability  $1/6$ .
  - Ex Event of rolling an even number, with probability  $1/2$ .  
Composed of basic events of rolling a 2, 4 or 6.
  - If A is event, write  $\Pr[A]=y$ . E.g.  $\Pr[\text{roll a 2}]=1/6$
- Two events A, B are independent if  $\Pr[A \wedge B] = \Pr[A] * \Pr[B]$ .
  - Ex Events A="2 on first roll" and B="3 on second roll" are independent, because  $\Pr[A \wedge B] = 1/36 = \Pr[A] * \Pr[B] = 1/6 * 1/6$ .
  - Ex Events A="2 on first roll" and B="the two rolls sum to 5" are not independent, because  $\Pr[A \wedge B] = 1/36 \neq \Pr[A] * \Pr[B] = 1/6 * 4/36$ .





# Probability review

## ■ Random variables

- A variable which takes values with certain probabilities. The probabilities sum to 1.
- **Ex**  $X$  = value from roll of dice. Values are  $\{1,2,3,4,5,6\}$ , each with probability  $1/6$ .
- **Ex**  $Y$  = number of heads in 4 flips of fair coin. Values are  $\{0,1,2,3,4\}$ , with probabilities  $\{1/16, 4/16, 6/16, 4/16, 1/16\}$ .
- **Ex**  $Z$  = number of flips of fair coin till first head. Values are  $\{1,2,3,\dots\}$ , with probabilities  $\{1/2, 1/4, 1/8, \dots\}$ .
- We write  $\Pr[X=x]=y$ , e.g.  $\Pr[Z=3]=1/8$ .



# Probability review

- Expectation of random variable  $X$

- $E[X] = \sum_x x \cdot \Pr[X=x]$ .

- The average value of  $X$ , over many trials.

- **Ex**  $X$  = number of heads in 4 flips.

- $E[X] = 0 \cdot 1/16 + 1 \cdot 4/16 + 2 \cdot 6/16 + 3 \cdot 4/16 + 4 \cdot 1/16 = 2.$

- If you flip a coin 4 times, for 1000 times, on average you see 2 heads per 4 flips.

- An indicator variable  $X$  for an event  $E$  is a random variable that's 1 if  $E$  occurs, and 0 otherwise.

- If event  $E$  has probability  $p$  of occurring, and  $X$  is  $E$ 's indicator variable, then  $E[X] = p$ .

- Because  $E[X] = \Pr[E \text{ occurs}] \cdot 1 + \Pr[E \text{ doesn't occur}] \cdot 0 = p$ .

- This is a convenient fact we'll frequently use.

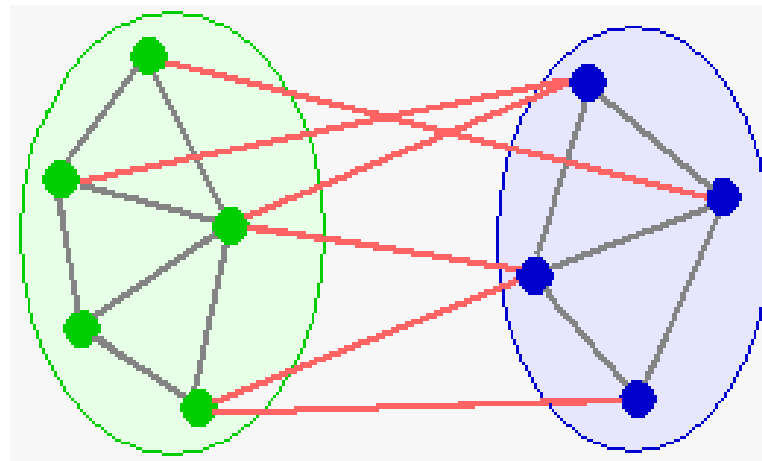
# Probability review

## ■ Linearity of expectations

- Given random variables  $X, Y$ ,  $E[X+Y]=E[X]+E[Y]$ .
- Extends to any number of random variables, e.g.  $E[X+Y+Z]=E[X]+E[Y]+E[Z]$ .
- The random variables do not have to be independent.
- Very useful property!
- **Ex** Let  $X$ =number of heads in 100 coin flips. Calculate  $E[X]$ .
  - **Direct method**:  $1*\text{Pr}[1 \text{ head}]+2*\text{Pr}[2 \text{ heads}]+\dots+100*\text{Pr}[100 \text{ heads}]$ , a very complicated sum.
  - **Linearity method**:  $X=X_1+X_2+\dots+X_{100}$ , where  $X_i$ =number of heads on  $i$ 'th flip.
  - $E[X_i]=0*\text{Pr}[0 \text{ heads}]+1*\text{Pr}[1 \text{ head}]=1/2$ .
  - $E[X]=E[X_1]+\dots+E[X_{100}]=100/2=50$ .

# Problem 1: Max-Cut

- We studied the Min-Cut problem, which is closely related to finding the max flow in a network.
- Max-Cut is the opposite of Min-Cut.
- Given a graph  $G$ , split vertices into two sides to maximize the number of edges between the sides.






# Max-Cut

- Unlike Min-Cut, **Max-Cut is NP-complete**.
- We'll give a very simple randomized Monte Carlo 2-approximation algorithm.
  - Monte Carlo means the algorithm sometimes returns the wrong answer, i.e. a cut that's not a 2-approximation.
  - Monte Carlo also means the algorithm always runs in a fixed amount of time.
- ❖ Put each node in a random side with probability  $\frac{1}{2}$ .

# Correctness

- **Lemma** In a graph with  $e$  edges, the algorithm produces a cut with expected size  $e/2$ .
- **Proof** Let  $X$  be a random variable equal to the size of the cut. We want to bound  $E[X]$ .
  - For each edge  $e$ , let  $X_e$  be the indicator variable of whether  $e$  is in the cut.
    - I.e.  $X_e=1$  if  $e$  is in the cut and 0 otherwise.
  - So  $X = \sum_e X_e$ .
  - Given an edge  $e=(i,j)$ ,  $e$  is in the cut if  $i$  and  $j$  are on different sides.
  - So  $\Pr[e \text{ in cut}] = \Pr[(i \text{ in } L) \wedge (j \text{ in } R)] + \Pr[(j \text{ in } L) \wedge (i \text{ in } R)] = 1/4 + 1/4 = 1/2$ .
  - So  $E[X_e] = 1/2$ .
  - So  $E[X] = e/2$  by linearity of expectations.



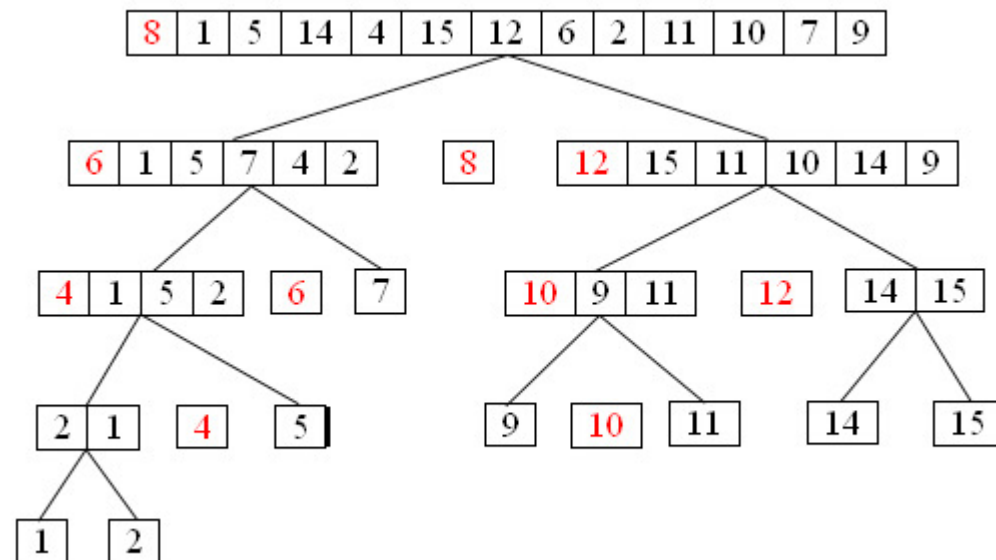


# Correctness

- Since a cut can have at most  $e$  edges, the  $e/2$  edges the algorithm outputs in expectation is a 2 approximation.
- Note that we only bounded expected size of the algorithm's cut.
  - In any particular execution, the algorithm can output a cut that's smaller or larger than  $e/2$ .
    - On average, the cut has size  $e/2$ .
  - It's possible to bound the probability the algorithm outputs a cut significantly smaller than  $e/2$ , but we won't do this.

# Problem 2: Quicksort


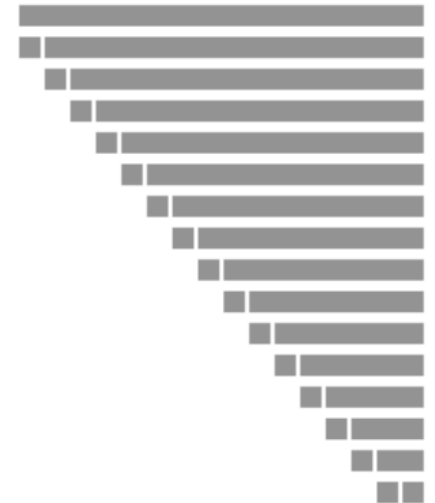
- Recall the Quicksort algorithm.
  - ❖ Pick a pivot element  $s$ .
  - ❖ Partition the elements into two sets, those less than  $s$  and those more than  $s$ .
  - ❖ Recursively Quicksort the two sets.





# Complexity of Quicksort

- Let  $T(n)$  be the time to Quicksort  $n$  numbers.
- $T(n)$  is small in practice.
- But in the worst case,  $T(n)=O(n^2)$ .
  - Occurs with very uneven splits. I.e. the rank of the pivot is very small or large.
  - **Ex** If pivot is smallest element, then  $T(n)=T(1)+T(n-1)+n-1$ . This solves to  $T(n)=O(n^2)$ .
    - $T(1)$  and  $T(n-1)$  to recursively sort each side,  $n-1$  to partition the elements wrt the pivot.
- As long as the pivot is near the middle, Quicksort takes  $O(n \log n)$  time.
  - **Ex** If the pivot is always in the middle half,  $[n/4, 3n/4]$ , then  $T(n) \leq T(n/4)+T(3n/4)+n-1$ , which solves to  $O(n \log n)$ .





# Randomized Quicksort

- Quicksort is only slow if we keep picking very small or large pivots.
- Let's pick the pivot at random. Intuitively, we shouldn't be unlucky and always pick small or large pivots.
- ❖ Pick a **random pivot** element  $s$ .
- ❖ Partition the elements into two sets, those less than  $s$  and those more than  $s$ .
- ❖ Recursively RQuicksort the two sets.



# Complexity of RQuicksort

- Let  $R(n)$  be the expected time to RQuicksort  $n$  numbers.
- With probability  $1/n$ , the pivot has rank 1 (is smallest element), in which case  $R(n)=R(1)+R(n-1)+n-1$ .
- With probability  $1/n$ , the pivot has rank 2, and  $R(n)=R(2)+R(n-2)+n-1$ .
- ...
- With probability  $1/n$ , the pivot has rank  $k$ , and  $R(n)=R(k)+R(n-k)+n-1$ .
- Putting these together, we have
$$R(n) = 1/n \cdot (R(1)+R(n-1)+R(2)+R(n-2)+\dots+R(n-1)+R(1) + n \cdot 1/n \cdot (n-1) =$$
$$2/n \cdot \sum_k R(k) + n-1.$$

# Complexity of RQuicksort

- We solve the recurrence for  $R(n)$  using the substitution method. We guess  $R(n) \leq an \log n + b$  for some constants  $a, b > 0$  to be determined.
- We first need the following lemma.

**Lemma 1.1**  $\sum_{k=1}^{n-1} k \log k \leq \frac{1}{2}n^2 \log n - \frac{1}{8}n^2$ .

**Proof.**

$$\begin{aligned} \sum_{k=1}^{n-1} k \log k &= \sum_{k=1}^{\lceil n/2 \rceil - 1} k \log k + \sum_{k=\lceil n/2 \rceil}^{n-1} k \log k \\ &\leq (\log n - 1) \sum_{k=1}^{\lceil n/2 \rceil - 1} k + \log n \sum_{k=\lceil n/2 \rceil}^{n-1} k \\ &= \log n \sum_{k=1}^{n-1} k - \sum_{k=1}^{\lceil n/2 \rceil - 1} k \\ &\leq \frac{1}{2}n(n-1) \log n - \frac{1}{2}\left(\frac{n}{2} - 1\right)\frac{n}{2} \\ &\leq \frac{1}{2}n^2 \log n - \frac{1}{8}n^2 \end{aligned}$$



# Complexity of RQuicksort

- Now we can solve for  $R(n)$ .

$$\begin{aligned} R(n) &= \frac{2}{n} \sum_{k=1}^{n-1} R(k) + \Theta(n) \\ &\leq \frac{2}{n} \sum_{k=1}^{n-1} (ak \log k + b) + \Theta(n) \\ &= \frac{2a}{n} \sum_{k=1}^{n-1} k \log k + \frac{2b(n-1)}{n} + \Theta(n) \\ &\leq \frac{2a}{n} \left( \frac{1}{2} n^2 \log n - \frac{1}{8} n^2 \right) + \frac{2b}{n} (n-1) + \Theta(n) \\ &\leq an \log n - \frac{a}{4} n + 2b + \Theta(n) \\ &= an \log n + b + (\Theta(n) + b - \frac{a}{4} n) \\ &\leq an \log n + b \end{aligned}$$

by choosing  $a$  so that  $\frac{a}{4}n > \Theta(n) + b$ .




# Hash tables

- A hash table is a randomized data structure to efficiently implement a dictionary.
- Supports find, insert, and delete operations all in expected  $O(1)$  time.
  - But in the worst case, all operations are  $O(n)$ .
  - The worst case is provably very unlikely to occur.
- A hash table does not support efficient min / max or predecessor / successor functions.
  - All these take  $O(n)$  time on average.
- A practical, efficient alternative to binary search trees if only find, insert and delete needed.

# Direct addressing


- Suppose we want to store (key, value) pairs, where keys come from a finite universe  $U = \{0, 1, \dots, m-1\}$ .
- Use an array of size  $m$ .
  - `insert(k, v)` Store  $v$  in array position  $k$ .
  - `find(k)` Return the value in array position  $k$ .
  - `delete(k)` Clear the value in array position  $k$ .
- All operations take  $O(1)$  time.
- The problem is, if  $m$  is large, then we need to use a lot of memory.
  - Uses  $O(|U|)$  space.
  - **Ex** For 32 bit keys, need 16 GB memory. For 64 bit keys, more memory than in world.
  - **Ex** What about string based keys?
- Also, if only need to store few values, lots of space wasted.



Source: Introduction to Algorithms, Cormen et al

# Hash table

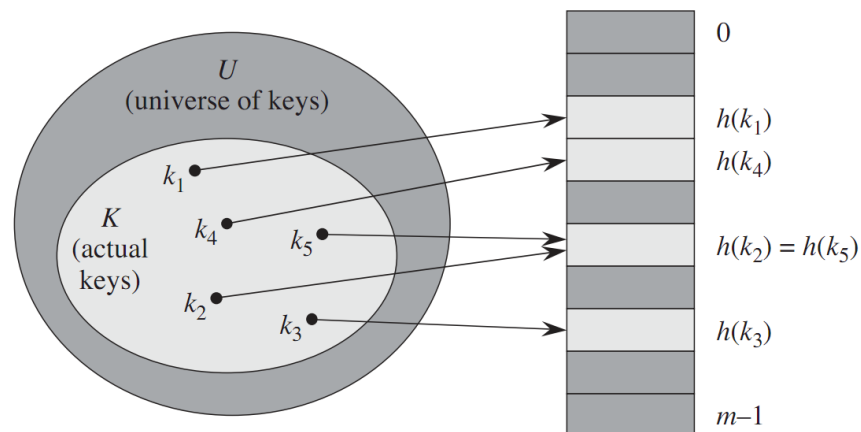
- Similar to direct addressing, but uses much less space.
- **Idea** Instead of storing directly at key's location, convert key to much smaller value, and store at this location.
- A hash table consists of the following.
  - A universe  $U$  of keys.
  - An array of  $T$  of size  $m$ .
  - A hashing function  $h:U \rightarrow \{0,1,\dots,m-1\}$ .
- We'll talk later about how to pick good hash functions.
- **insert( $k, v$ )** Hash key to  $h(k)$ . Store  $v$  in  $T[h(k)]$ .
- **find( $k$ )** Return the value in  $T[h(k)]$
- **delete( $k$ )** Delete the value in  $T[h(k)]$
- Assuming  $h(k)$  takes  $O(1)$  time to compute, all ops still take  $O(1)$  time. Uses  $O(m)$  space.
- If  $m \ll |U|$ , then hashing uses much less space than direct addressing.
- However, our current scheme doesn't quite work, due to collisions.






# Collisions

- We store a key at array position  $h(k)$ .
- But what if two keys hash to the same location, i.e.  $k_1 \neq k_2$ , but  $h(k_1) = h(k_2)$ ?
  - This is called a collision.
- Collisions are unavoidable when  $|U| > m$ .
  - By Pigeonhole Principle, must exist at least two different keys in  $U$  that hash to same value.
- Two basic ways to deal with collisions, closed and open addressing.




# Closed addressing

- In closed addressing, every entry in hash table points to a linked list.
  - Keys that hash to the same location get added to the linked list.
  - For simplicity, we'll ignore values from now on and only focus on keys.
- **insert(k)** Add  $k$  to the linked list in  $T[h(k)]$ .
- **find(k)** Search the linked list in  $T[h(k)]$  for  $k$ .
- **delete(k)** Delete  $k$  from the linked list in  $T[h(k)]$ .
- Suppose the longest list has length  $\hat{n}$ , and average length list is  $\bar{n}$ .
  - Each operation takes worst case  $O(\hat{n})$  time.
  - An operation on a random key takes  $O(\bar{n})$  time.



# Load factor

- The key to making closed addressing hashing fast is to make sure list lengths aren't too long.
- For this, we want the hash function to appear random.
  - Assume that any key is uniformly likely to be hashed to any table location.
- Suppose the hash table contains  $n$  items, and has size  $m$ .
- Then under the uniform hashing assumption, each table location has on average  $n/m$  keys.
  - Call  $\alpha = n/m$  the load factor.
- So the average time for each operation is  $O(\alpha)$ .
- However, even with uniform hashing, in the worst case, all keys can hash to the same location. So the worst case performance is  $O(n)$ .





# Picking a hash function

- We saw that we want hash functions to hash keys to “random” locations.
  - However, note that each hash function is itself a deterministic function, i.e.  $h(k)$  always has the same value.
    - If  $h(k)$  can produce different values, we can’t find key  $k$  in the hash table anymore.
- It’s hard to find such random hash functions, since we don’t assume anything about the distribution of input keys.
  - Ex For any hash function, there are always  $\geq |U|/m$  keys from the universe hashing to the same location. So if the input is exactly this set, and  $|U|/m \geq n$ , then all ops take  $O(n)$  time.
- In practice, we use a number of heuristic functions.



# Heuristic hash functions

- Assume the keys are natural numbers.
  - Convert other data types to numbers.
  - **Ex** To convert ASCII string to natural number, treat the string as a radix 128 number. E.g. “pt”  
 $\rightarrow (112 \cdot 128) + 116 = 14452$ .
- **Division method**  $h(k) = k \bmod m$ 
  - Often choose  $m$  a prime number not too close to a power of 2.
- **Multiplication method**  $h(k) = \lfloor m (k A \bmod 1) \rfloor$ , where  $A$  is some constant.
  - Knuth's suggestion is  $A = \frac{\sqrt{5}-1}{2} \approx 0.618034 \dots$



# Universal hashing

- As we said, regardless of the hash function, an adversary can choose a set of  $n$  inputs to make all operations  $O(n)$  time.
- Universal hashing overcomes this using randomization.
  - No matter what the  $n$  input keys are, every operation takes  $O(n/m)$  time in expectation, for a size  $m$  hash table.
  - Note  $O(n/m)$  time is optimal.
- Instead of using a fixed hash function, universal hashing uses a random hash function, chosen from some set of functions  $H$ .
- Say  $H$  is a universal hash family if for any keys  $x \neq y$

$$\Pr_{h \in H} [h(x) = h(y)] = 1/m$$

- So if we randomly choose a hash function from  $H$  and use it to hash any keys  $x, y$ , they have  $1/m$  probability of colliding.
- Note the hash functions in  $H$  are not random. However, we choose which function to use from  $H$  randomly.

# Universal hashing

- **Thm** Let  $H$  be a universal hash family. Let  $S$  be a set of  $n$  keys, and let  $x \in S$ . If  $h \in H$  is chosen at random, then the expected number of  $y \in S$  s.t.  $h(x) = h(y)$  is  $n/m$ .
- **Proof** Say  $S = \{x_1, \dots, x_n\}$ .
  - Let  $X$  be a random variable equal to the number of  $y \in S$  s.t.  $h(x) = h(y)$ .
  - Let  $X_i = 1$  if  $h(x_i) = h(x)$  and 0 otherwise.
  - $E[X_i] = \Pr_{h \in H} [h(x_i) = h(x)] \times 1 + \Pr_{h \in H} [h(x_i) \neq h(x)] \times 0 = 1/m$ .
    - First equality follows by universal hashing property.
  - $E[X] = E[X_1] + \dots + E[X_n] = n/m$ .

# Constructing universal hash family 1

- Choose a prime number  $p$  such that  $p > m$ , and  $p >$  all keys.
- Let  $h_{ab}(k) = ((ak + b) \bmod p) \bmod m$ .
- Let  $H_{pm} = \{h_{ab} \mid a \in \{1, 2, \dots, p-1\}, b \in \{0, 1, \dots, p-1\}\}$ .
- **Thm**  $H_{pm}$  is a universal hash family.
- **Proof** Let  $x, y < p$  be two different keys. For a given  $h_{ab}$  let
$$r = (ax + b) \bmod p, \quad s = (ay + b) \bmod p$$
- We have  $r \neq s$ , because  $r - s \equiv a(x - y) \bmod p \neq 0$ , since neither  $a$  nor  $x - y$  divide  $p$ .
- Also, each pair  $(a, b)$  leads to a different pair  $(r, s)$ , since
$$a = ((r - s)(x - y)^{-1} \bmod p), \quad b = (r - ax) \bmod p$$
  - Here,  $(x - y)^{-1} \bmod p$  is the unique multiplicative inverse of  $x - y$  in  $\mathbb{Z}_p^*$ .



# Constructing universal hash family 2

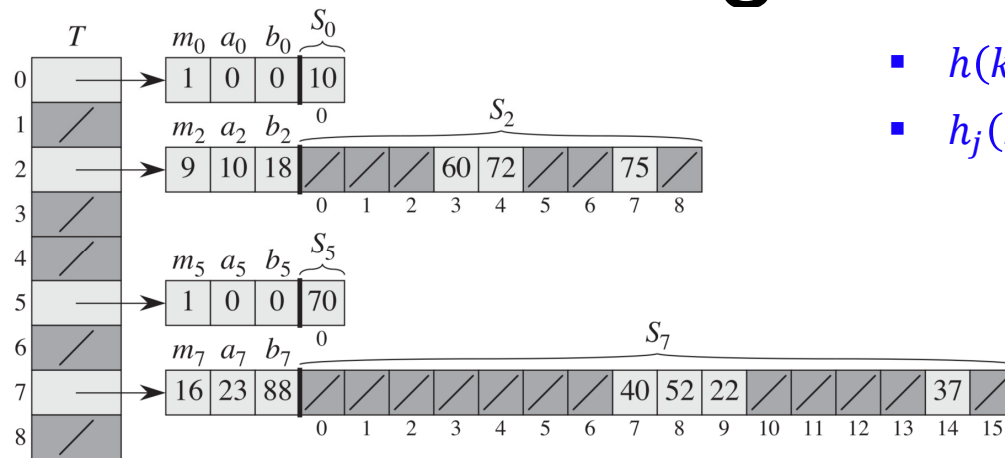
- Since there are  $p(p - 1)$  pairs  $(a, b)$  and  $p(p - 1)$  pairs  $(r, s)$  with  $r \neq s$ , then a random  $(a, b)$  produces a random  $(r, s)$ .
- The probability  $x$  and  $y$  collide equals the probability  $r \equiv s \pmod{m}$ .
- For fixed  $r$ , number of  $s \neq r$  s.t.  $r \equiv s \pmod{m}$  is  $(p - 1)/m$ .
- So for each  $r$  and random  $s \neq r$ , probability that  $r \equiv s \pmod{m}$  is  $((p - 1)/m)/(p - 1) = 1/m$ .
- So  $\Pr_{h_{ab} \in H_{pm}} [h_{ab}(x) = h_{ab}(y)] = 1/m$  and  $H_{pm}$  is universal.



# Perfect hashing

- The hashing methods we've seen can ensure  $O(1)$  expected performance, but are  $O(n)$  in the worst case due to collisions.
- However, if we have a fixed set of keys, perfect hashing can ensure **no collisions** at all.
  - Perfect hashing maintains a static set, and allows  $\text{find}(k)$  and  $\text{delete}(k)$  in  $O(1)$  time.
  - It doesn't support  $\text{insert}(k)$ .
- **Ex** The fixed set of keys may represent the file names on a non-writable DVD.

# Perfect hashing



- $h(k) = ((3k + 42) \bmod 101) \bmod 9$
- $h_j(k) = ((a_j k + b_j) \bmod 101) \bmod m_j$

- Suppose we want to store  $n$  items with no collisions.
- Perfect hashing uses two levels of universal hashing.
  - The first layer hash table has size  $m = n$ .
  - Use first layer hash function  $h$  to hash key to a location in  $T$ .
  - Each location  $j$  in  $T$  points to a hash table  $S_j$  with hash function  $h_j$ .
  - If  $n_j$  keys hash to location  $j$ , the size of  $S_j$  is  $m_j = n_j^2$ .
- We'll ensure there are no collisions in the secondary hash tables  $S_1, \dots, S_m$ .
  - So all operations take worst case  $O(1)$  time.
- Overall the space use is  $O(m + \sum_{j=1}^m n_j^2)$ .
  - We'll show this is  $O(n) = O(m)$ .
  - So perfect hashing uses same amount of space as normal hashing.

# Avoiding collisions

- **Lemma** Suppose we store  $n$  keys in a hash table of size  $m = n^2$  using universal hashing. Then with probability  $\geq 1/2$  there are no collision.
- **Proof** There are  $\binom{n}{2}$  pairs of keys that can collide.
  - Each collision occurs with probability  $1/m = 1/n^2$ , by universal hashing.
  - So the expected number of collisions is  $\frac{\binom{n}{2}}{n^2} \leq \frac{1}{2}$ .
  - By Markov's inequality the  $\Pr[\# \text{ collisions} \geq 1] \leq E[\# \text{ collisions}] \leq 1/2$ .
- When building each hash table  $S_j$ , there's  $< 1/2$  probability of having any collisions.
  - If collisions occur, pick another random hash function from the universal family and try again.
  - In expectation, we try twice before finding a hash function causing no collisions.

# Space complexity

- **Lemma** Suppose we store  $n$  keys in a hash table of size  $m=n$ . Then the secondary hash tables use space  $E[\sum_{j=0}^{m-1} n_j^2] \leq 2n$ , where  $n_j$  is the number of keys hashing to location  $j$ .
- **Proof**  $E[\sum_{j=0}^{m-1} n_j^2] = E[\sum_{j=0}^{m-1} (n_j + 2 \binom{n_j}{2})] = E[\sum_{j=0}^{m-1} n_j] + 2 E[\sum_{j=0}^{m-1} \binom{n_j}{2}]$
- $\sum_{j=0}^{m-1} \binom{n_j}{2}$  is the total number of pairs of hash keys which collide in the first level hash table.
  - By universal hashing, this equals  $\binom{n}{2} \frac{1}{m} = \frac{n-1}{2}$ .
- $E[\sum_{j=0}^{m-1} n_j] = n$ .
- So  $E[\sum_{j=0}^{m-1} n_j^2] \leq n + \frac{2(n-1)}{2} < 2n$ .