

Approximation algorithms 3

Scheduling, Knapsack


CS240

Spring 2022

Rui Fan

The knapsack problem

- We have a set of items, each having a weight and a value.
- We have a knapsack that can carry up to W amount of weight.
- We want to put items in the knapsack to maximize the total value, but not exceed the weight limit.
- **Ex** Items 3 and 4 are the highest value items with weight ≤ 11 .
- Assume all items have weight $\leq W$, i.e. any single item fits in knapsack.



$W = 11$

Item	Value	Weight
1	1	1
2	6	2
3	18	5
4	22	6
5	28	7

A dynamic program for knapsack

- Let $\text{OPT}(i,v)$ = minimum weight of a subset of items $1,\dots,i$ that has value $\geq v$.
- If optimal solution uses item i .
 - Then we pay w_i weight for item i , and need to achieve value $\geq v-v_i$ using items $1,\dots,i-1$ using min weight.
 - So $\text{OPT}(i,v)=w_i+\text{OPT}(i-1,v-v_i)$.
- If optimal solution doesn't use item i .
 - Then we need to achieve value $\geq v$ using items $1,\dots,i-1$.
 - So $\text{OPT}(i,v)=\text{OPT}(i-1,v)$.
- Choose the case that gives smaller weight.
- $$\text{OPT}(i,v) = \begin{array}{ll} 0 & \text{if } v=0 \\ \infty & \text{if } i=0, v>0 \\ \min(\text{OPT}(i-1,v), w_i+\text{OPT}(i-1,v-v_i)) & \text{otherwise} \end{array}$$



Running time of dynamic program

- Say there are n items, and the largest value of any item is v^* .
- The max value we can pack into the knapsack is nv^* , where v^* is the largest v value.
- Solve all subproblems of the form $\text{OPT}(i,v)$, where $i \leq n$ and $v \leq nv^*$.
 - This is a total of $O(n^2v^*)$ subproblems.
- The solution to Knapsack is the max value V that can be packed with weight $\leq W$.
- Having solved all the subproblems, we can find V by finding the subproblem with the largest value that has optimum weight $\leq W$.
 - $V = \max_{v \leq nv^*} \text{OPT}(n,v) \leq W$.
- So solving Knapsack takes total time $O(n^2v^*)$.



Running time of dynamic program

- The DP gives an optimal solution to Knapsack and takes $O(n^2v^*)$ time. Have we found a polytime algorithm for an NP-complete problem?
- No. The problem size is $O(n \log(v^*))$, because it takes $\log(v^*)$ bits to express each item's value. But $O(n^2v^*)$ is not polynomial in $n \log(v^*)$.
- To make this DP fast, we have to make the largest value small.



PTAS

- Let $\varepsilon > 0$ be any number. We'll give a $(1+\varepsilon)$ -approximation for knapsack.
- By setting ε sufficiently small, we can get as good an approximation as we want!
 - This type of algorithm is called a polynomial time approximation scheme, or PTAS.
- Contrast this with earlier algs we studied, which had worse approx ratios, e.g. 2 or $\log n$.
- But the running time will be $O(n^3/\varepsilon)$. Hence we can't set $\varepsilon=0$ get the optimal solution.
- We're trading accuracy for time. The more accurate (smaller ε), the more time the algorithm takes.

Main idea: rounding

- Since we only need an approximate solution, we can change the values of the items a little (round the values) and not affect the solution much.
- We scale and round the original values to make them small.
- The previous DP took $O(n^2v^*)$ time. So if the rounded values are small, this DP is fast.

W = 11						W = 11		
Item	Value	Weight				Item	Value	Weight
1	134,221	1				1	2	1
2	656,342	2				2	7	2
3	1,810,013	5				3	19	5
4	22,217,800	6				4	223	6
5	28,343,199	7				5	284	7



Rounding

- Let $\epsilon > 0$ be the precision we want.
- Set $\theta = \epsilon v^*/2n$ to be a scaling factor.
 - v^* is the largest value of any item.
- Scale all values down by θ then round up.
 - $v' = \lceil v/\theta \rceil$.
- Make a problem where each value v_i is replaced by v'_i .
 - Call this the scaled rounded problem.
- Let v^\wedge be max value in the scaled rounded problem. Then $v^\wedge = \lceil v^*/\theta \rceil = \lceil v^*/(\epsilon v^*/2n) \rceil = \lceil 2n/\epsilon \rceil$.
- Running time of DP on scaled rounded problem is $O(n^2 v^\wedge) = O(n^3/\epsilon)$.



Solving the original problem

- Make another new problem in which each value v_i is replaced by $u_i = \lceil v_i/\theta \rceil * \theta$.
 - Call this the rounded problem.
 - We have $u_i \geq v_i$, and $u_i \leq v_i + \theta$.
- Note u values are equal to v' values multiplied by θ .
 - Thus, the optimal solution for the rounded problem and the scaled rounded problem are the same.
- We now have 3 problems, the original problem, the scaled rounded problem, and the rounded problem.
- Let S be the optimal solution to the scaled rounded problem, which we can find in time $O(n^3/\epsilon)$. S is also optimal for the rounded problem.
- We'll show S is a $1+\epsilon$ approximation for the original problem.

Correctness

- **Thm** Let S^* be the optimal solution to the original problem. Then $(1+\varepsilon) \sum_{i \in S} v_i \geq \sum_{i \in S^*} v_i$.
Hence S is a $(1+\varepsilon)$ -approximate solution.

- **Proof**

$$\sum_{i \in S^*} v_i \leq \sum_{i \in S^*} u_i$$

$$u_i \geq v_i$$

$$\leq \sum_{i \in S} u_i$$

S is opt soln for rounded problem

$$\leq \sum_{i \in S} (v_i + \theta)$$

$$u_i \leq v_i + \theta$$

$$\leq \sum_{i \in S} v_i + n \theta$$

$$|S| \leq n$$

Correctness

- Suppose item j has the largest value, so $v^* = v_j$. Then $n\theta = \frac{\varepsilon}{2} v_j \leq \frac{\varepsilon}{2} u_j \leq \frac{\varepsilon}{2} \sum_{i \in S} u_i$

□ Last inequality because item j itself is feasible solution, so opt solution S is no smaller.

- So $\sum_{i \in S} v_i \geq \sum_{i \in S} u_i - n\theta \geq \left(\frac{2}{\varepsilon} - 1\right) n\theta$, where first inequality comes from inequalities on last page.

- Assuming $\varepsilon \leq 1$, then $n\theta \leq \varepsilon \sum_{i \in S} v_i$

- Finally, we have

$$\sum_{i \in S^*} v_i \leq \sum_{i \in S} v_i + n\theta \leq \sum_{i \in S} v_i + \varepsilon \sum_{i \in S} v_i = (1 + \varepsilon) \sum_{i \in S} v_i$$



Summary

- We gave a DP for Knapsack.
- We scale and round to reduce number of different item values.
- Running the DP on the scaled rounded problem and using the solution for the original problem leads to an arbitrarily good approximation for Knapsack, a PTAS.
- There are PTAS's for a number of other problems.
 - Multiprocessor scheduling.
 - Bin packing.
 - Euclidean TSP.
- However, there are also many problems for which PTAS's do not exist, unless $P=NP$.

A decorative graphic on the left side of the slide consists of a large, light blue square that overlaps a smaller, slightly offset light blue square. Overlapping these is a series of smaller squares in various shades of blue, arranged in a stepped pattern. A solid dark blue horizontal bar extends from the right side of this pattern across the middle of the slide.

Lower Bounds



Upper and lower bounds

- What is the minimum resources (time, space, etc.) needed to solve a problem?
- Consider sorting n numbers.
 - Insertion sort takes $O(n^2)$ time.
 - This puts an upper bound of $O(n^2)$ on the time to sort n numbers.
 - Merge sort takes $O(n \log n)$ time.
 - This puts an upper bound of $O(n \log n)$ on the time to sort n numbers.
- We want to make the upper bound as low as possible, i.e. solve the problem faster.
- Suppose an algorithm A solves problem X in $f(n)$ time when input size is n .
 - Then $f(n)$ is an upper bound on the complexity of X .



Upper and lower bounds

- What about the least amount of time to solve X ?
- Suppose we know that any algorithm that solves X takes at least $g(n)$ time, when X has size n .
 - Then $g(n)$ is a lower bound on the complexity of X .
- If the lower bound $g(n)$ is large, it means problem X is hard to solve.
 - Ex NP-Hard problems are hard because they (probably) have super-polynomial lower bounds.
- To show a lower bound, we need to give a proof.
 - Usually we show if an algorithm takes too little time, it must sometimes produce the wrong answer.
- The lower bound for a problem depends on the computational model.
 - If a model has very powerful primitive operations, then algorithms can run faster, and the lower bound is smaller.
- If the complexity of an algorithm for problem X matches the lower bound for problem X , the algorithm is optimal, and the lower bound is tight.



Sorting

- How many comparisons are needed to sort n numbers?
- Upper bound: $O(n \log n)$ using merge sort.
- Lower bound: $\Omega(n \log n)$.
- To prove the lower bound, we first need a model for how a comparison-based sorting algorithm works.
 - This is called the decision tree model.
- The lower bound is not valid in other models.
 - If an algorithm can do things besides comparing two numbers, e.g. look at the digits of a number, it can sort faster than $\Omega(n \log n)$ time.
 - Lower bounds can be very sensitive to the computational model.




Decision trees

- In this model, in each step, algorithm can only compare a pair of numbers x , y .
- Based on result of the comparison, it decides next pair of numbers to compare.
 - So an execution of the algorithm is a sequence of comparisons, each comparison determined by result of previous comparison.
- When the algorithm terminates, it outputs a permutation representing the sorted order of the input.
- The complexity of the algorithm is the most number of comparisons it does before terminating.

Decision trees

- Model behavior of the algorithm by a binary tree.
 - Each internal node is a pair of number x,y to compare.
 - If $x \leq y$, go to left child. If $x > y$, go to right child.
 - Each leaf represents an output, and is labeled with a permutation representing the sorted order of the inputs.
- An execution is simply a path from root to a leaf.
 - At any node, the algorithm has obtained some info from the comparisons it's done.
 - It uses this info to decide the next comparison to do.
 - Eventually, it obtains enough info to generate an output.
- Complexity of algorithm is the length of the longest root-leaf path.





Lower bound for sorting

- Given n numbers as input, they can be in $n!$ different orders.
- Given an input order, algorithm must output that order.
 - So decision tree of algorithm must have a leaf labeled with that order.
 - So the decision tree has $\geq n!$ leaves.
- Say height of decision tree is h .
 - The complexity of the algorithm is h .
 - Since decision tree is binary, it has $\leq 2^h$ leaves.
- So $2^h \geq (\# \text{ leaves of dec tree}) \geq n!$, and so $h \geq \log_2(n!)$.
 - $\log_2(n!) = \log_2 n + \log_2(n-1) + \dots + \log_2 1 \geq \log_2 n + \log_2(n-1) + \dots + \log_2(n/2) \geq \frac{n}{2}(\log_2 n - 1) = \Omega(n \log n)$.
 - Can also use Stirling's approximation.
- So we proved the algorithm does $\Omega(n \log n)$ comparisons.